

Table of Contents

<i>1. Project Title, Authors, and Stakeholder</i>	<i>2</i>
<i>2. Preface.....</i>	<i>3</i>
<i>3. Introduction.....</i>	<i>4</i>
<i>4. Important Terms Definition & Introduction</i>	<i>6</i>
<i>5. Sequence Diagram & UML Class Diagram</i>	<i>10</i>
<i>6. Current Main Components & Status</i>	<i>11</i>
6a. Tigo's Town Component:	11
6b. Ricardo's Highway Component (Two main functions):.....	12
6c. Dongming's Highway Layout Algorithm Component:	13

1. Project Title, Authors, and Stakeholder

- a. Project Title: *Life-Long MAPF Distributed System*
- b. Team/Project number: 26
- c. Team name: MAPFer
- d. Team members:
 - Dongming Shen, USCID: 7170243072
 - Tiancheng Xu, USCID: 1522206865
 - Yi Jiang, USCID 5261852930
- e. Stakeholder: Christopher Leet
- f. TA: Marcelo Laser

2. Preface

This document provides a detailed process report (Until Oct. 30, 2022) for the Life-Long Distributed MAPF (multi-agent path finding) System project.

3. Introduction

Recent years have seen multi-robot systems increase in size and ubiquity. Warehouse operators such as Amazon routinely service large warehouse complexes with teams of thousands of robots. Entertainment companies such as Spaxels produce aerial spectacles consisting of thousands of drones. Each of these systems must solve the Multi-Agent Path Finding (MAPF) problem, the problem of finding collision-free paths for a group of agents operating in a known common environment. A practical solver must:

- (a) Produce near-optimal paths
- (b) Solve very large (multi-thousand robots) MAPF instances in real-time
- (c) Be capable of adjusting its solution if a robot completes its goal and is assigned a new goal, and
- (d) Be capable of adjusting its solution if a robot's movement fails.

To date, industrial multi-robot systems solve the MAPF problem using a traffic system-based approach. In a traffic system-based approach, the environment that agents operate in is divided into components such as roads and intersections, much like a city traffic system. Simple local rules govern the behavior of agents in each component. This approach is scalable and robust to new goal assignment and movement failure but is often far from optimal.

In recent years, academia has seen the development of search-based approaches to MAPF, such as conflict-based search. These approaches can be made optimal/bounded suboptimal but fail to scale and struggle to handle movement failure and goal reassignment.

A hybrid approach promises to combine the best of both systems. In this approach, open spaces controlled by a search-based MAPF algorithm are connected using a traffic-system-based

algorithm. Each search-based MAPF instance and components of connecting traffic system are run in a separate thread, providing a high degree of scalability. Only one component must replan should an agent's movement fail or goal change, providing robustness to these conditions.

This project involves key ideas from distributed systems, rule-based AI, and search-based AI hot topics in modern industry in the context of a real research problem with direct application to industry. The main goal of this project is to build a system that can solve MAPF problems involving over ~10000 agents with acceptable efficiency.

4. Important Terms Definition & Introduction

Highway:

The road system. Multiple highways will connect the towns. The agent needs to go through these highways to go from one Town to another. It will take agents who just leave a town and then plan the paths for them to arrive at their destination town where their goals are. Use a rule-based MAPF algorithm.

Highway Node:

Each grid/node in this highway system (these nodes compose the highways). Each of these nodes should either have exactly one out-going direction (w, e, n, s) or be an intersection (multiple out-going directions).

Town:

The map is split into different regions known as towns. Each Town runs independently to direct the agents to their goals. Each Town takes the new agents from the Highway through locks, runs EECBS to plan the paths for agents in it, and sends agents to the Highway with locks. Agents stay around in the Town while the Highway is crowded.

Town Node:

Each grid/node in a town. Town grid can be either lock('L'), obstacle('@'), or empty cell('.'). Each node is connected to another one in four directions: up, down, left, and right.

Lock Node:

Serves the buffer node between Town & Highway to avoid bottleneck crowding agents that negatively affect the system's throughput. For example, if many agents want to go from Town_1 to Town_2 through Highway simultaneously while Highway is narrow, then in Chris's previous system, agents might get stuck in Highway and, therefore, also stuck around at Town_1 to Highway's entrance. However, adding the "Lock" concept, most agents in Town_1 will not even get close to Highway's entrance with a high probability. This is because before an agent in Town_1 is assigned a Lock_out (to enter Highway from Town_1) and a Lock_in (to enter Town_2 from Highway), it will wander around in Town_1. In this way, an agent in Town_1 will not get stuck around the Highway entrance but will wait in other places until it gets "called" and can then go to Town_2 more efficiently with a much lower probability of getting stuck. This idea aims to improve the system's throughput.

Global Coordinator:

The Global Coordinator serves to 1. Get new tasks (goals) for the free agents in Town, and assign Locks (if available) to the agents, or set the agent's property so that it will wander around in its Town; 2. Handle messages (thread communications) between the Town system and the Highway system; 3. Free locks when an agent finishes using it.

Agent:

Each agent has a current_location and a goal_location (parsed by the Global Coordinator from Tasks). Usually, current_location and goal_location are in different towns (for example, Town_1

and Town_2). Then after assigning the agent's goal_location, the Global Coordinator should also check if there is a (Lock_out (in Town_1), Lock_in (in Town_2)) pair, where both Locks are free and connected by Highway. If it exists, the Global Coordinator should assign Lock_out and Lock_in to the agent and lock these Locks to prevent conflicts. If it does not exist, the Global Coordinator should assign the agent a random goal in Town_1 to make the agent wander around. In rare cases, if current_location and goal_location are in the same Town, then the Global Coordinator should not worry about the Locks but will simply assign goal_location to the agent for the Town to route the agent to its goal directly.

Routing Agent:

(We will focus on the general case where an agent has current_location in Town_1, goal_location in Town_2, and is assigned Lock_out in Town_1 and Lock_in in Town_2. This is because other cases are only sub-problems of this general case.)

Here are the steps to route the agent from its current_location to its goal_location:

1. Town_1 will route the agent from its current_location to any empty adjacent node Adj_Lock_1 of Lock_1 in Town_1.
2. Upon arriving at the Adj_Lock_1, Town_1 tells Global this information; Global then tells Highway this information; Highway blindly transfers the agent from Adj_Lock_1 to Lock_1 (which should be empty and is therefore safe).
3. Highway will route the agent from Lock_1 to Lock_2. Upon leaving Lock_1, Highway will tell Global this information; then Global will free Lock_1.

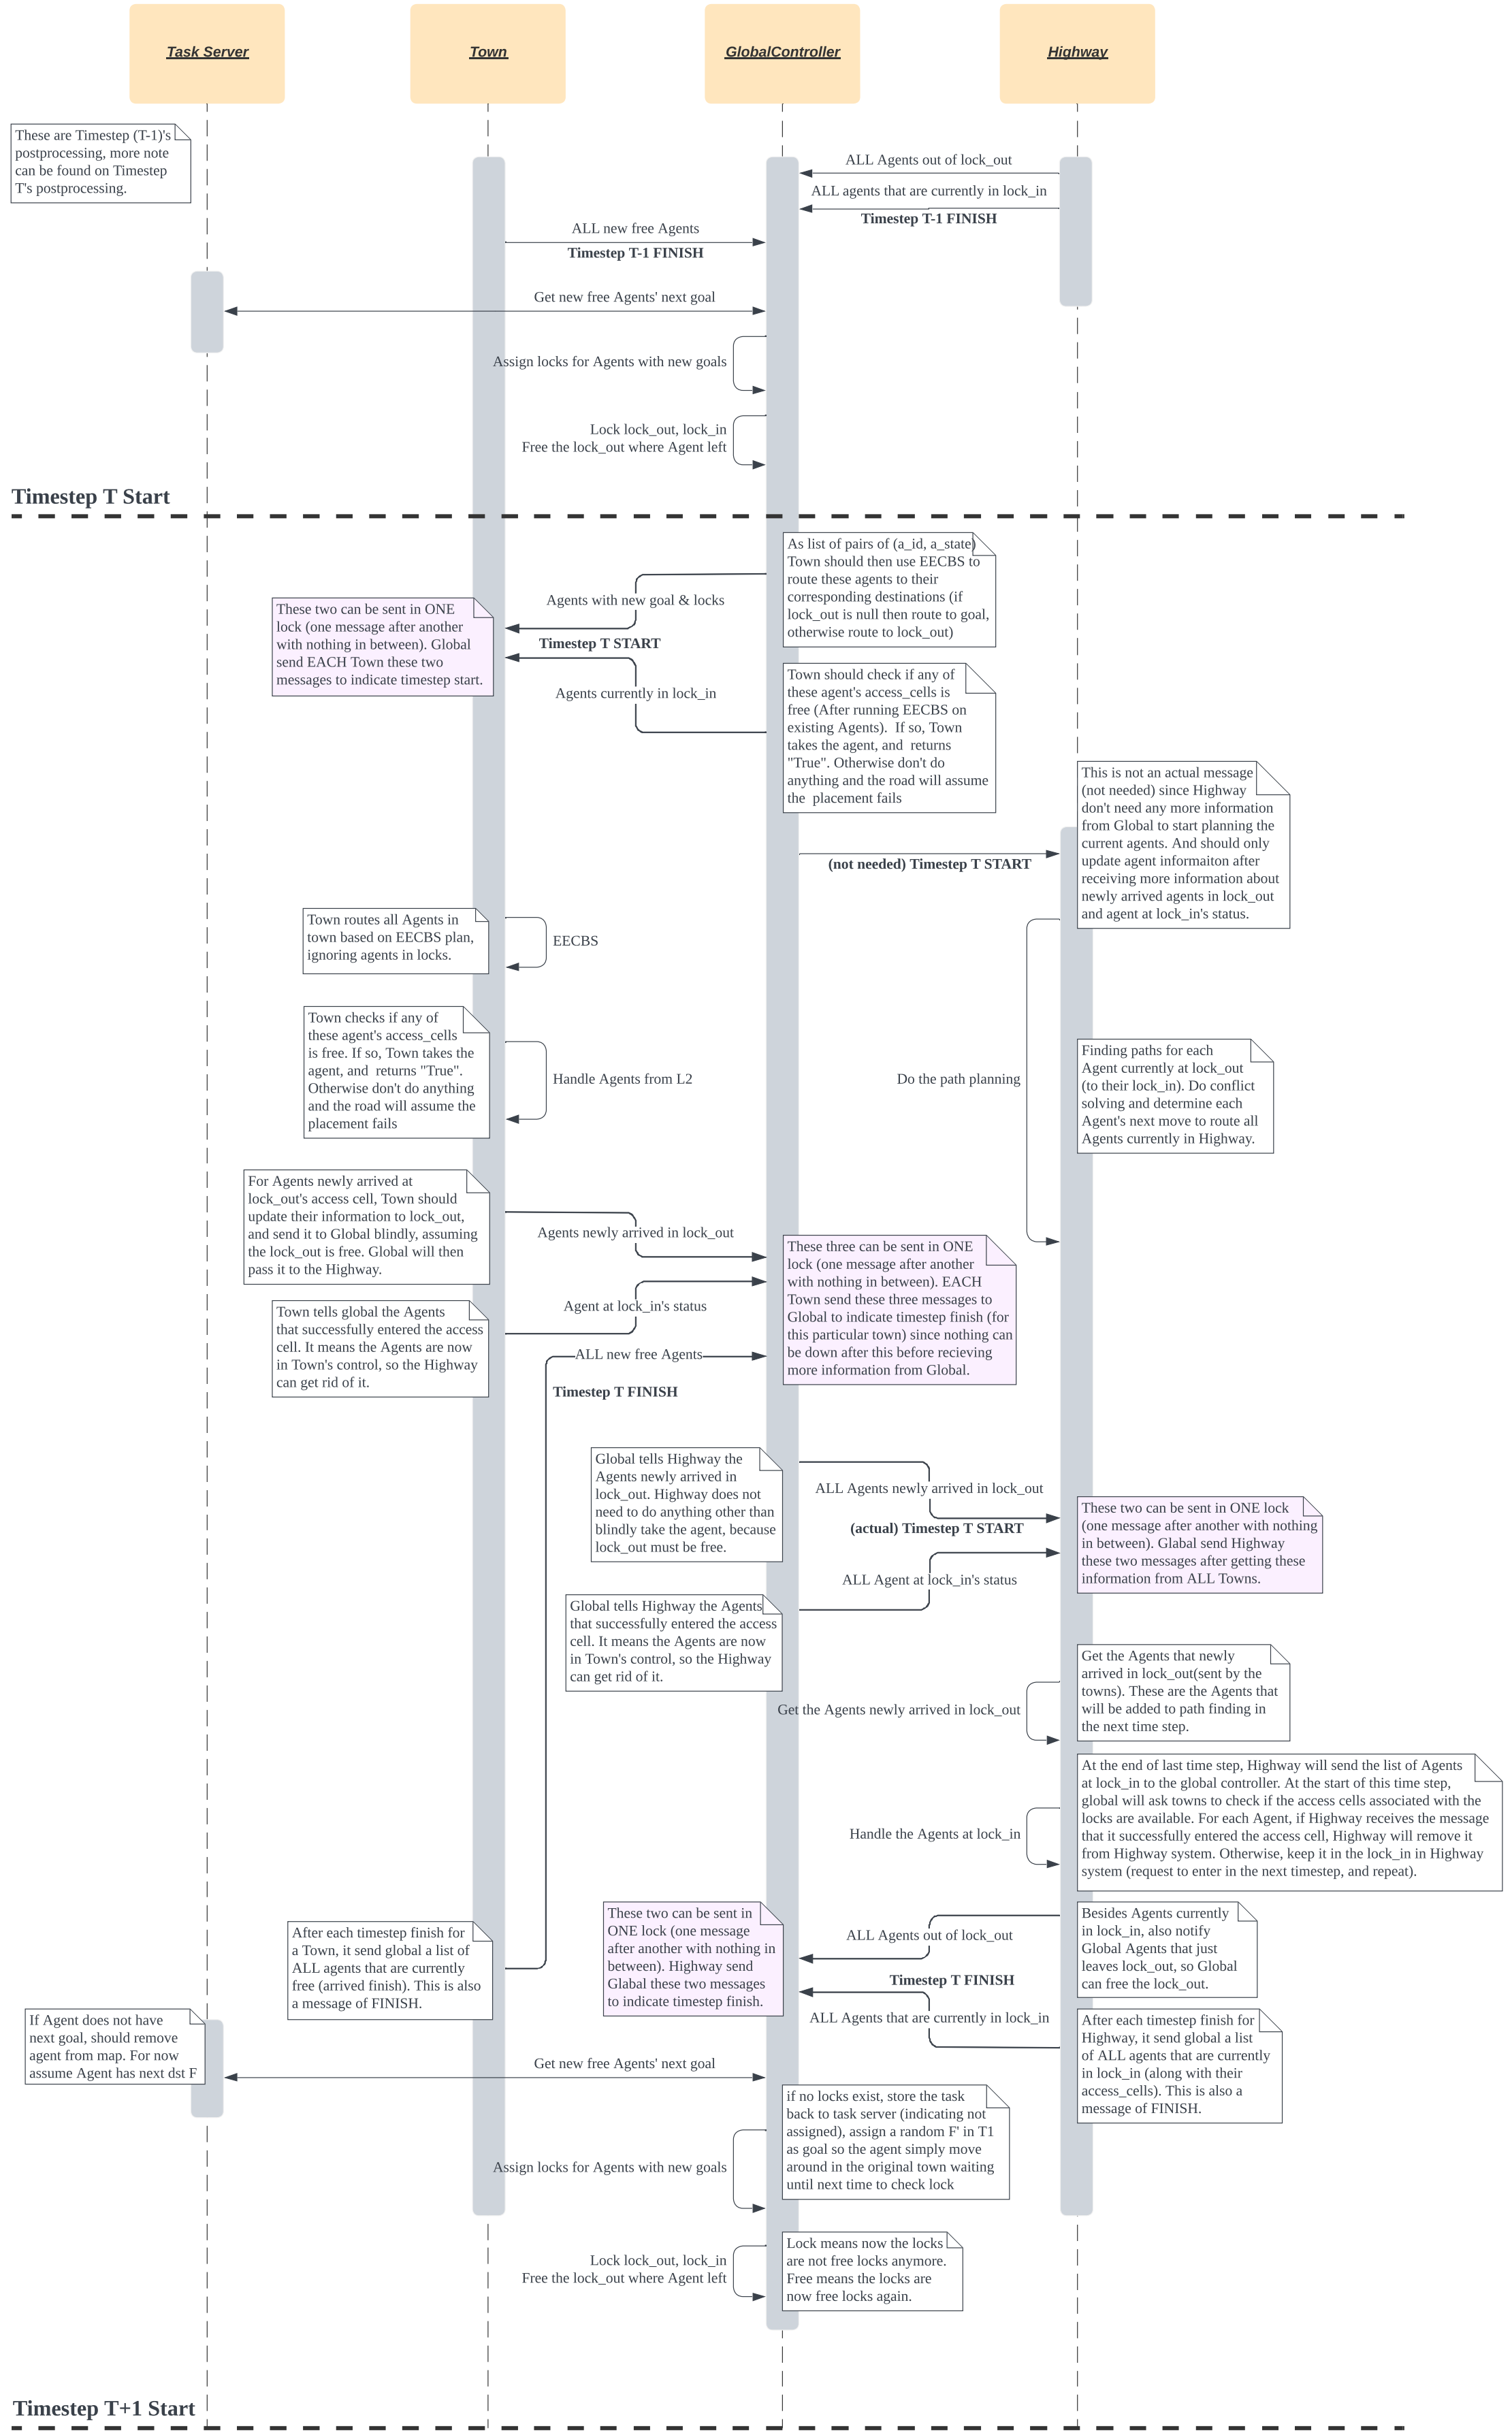
4. Upon arriving at Lock_2, Highway tells Global this information, Global then tells Town_2 this information. Town_2 checks if any adjacent node of Lock_2 in Town_2 is empty: if so, it transfers the agent from Lock_2 to this empty Adj_Lock_2 and tells Global the information; Global then tells Highway this information, so Highway will remove the agent from Lock_2; Global will also free Lock_2; if not, nothing happens, the agent will stay in Lock_2 and keep requesting every timestep until getting transferred into Town_2.
5. Finally, Town_2 will route the agent from Adj_Lock_2 to goal_location. When it arrives, Town_2 tells Global that the agent has reached the goal and is now free.

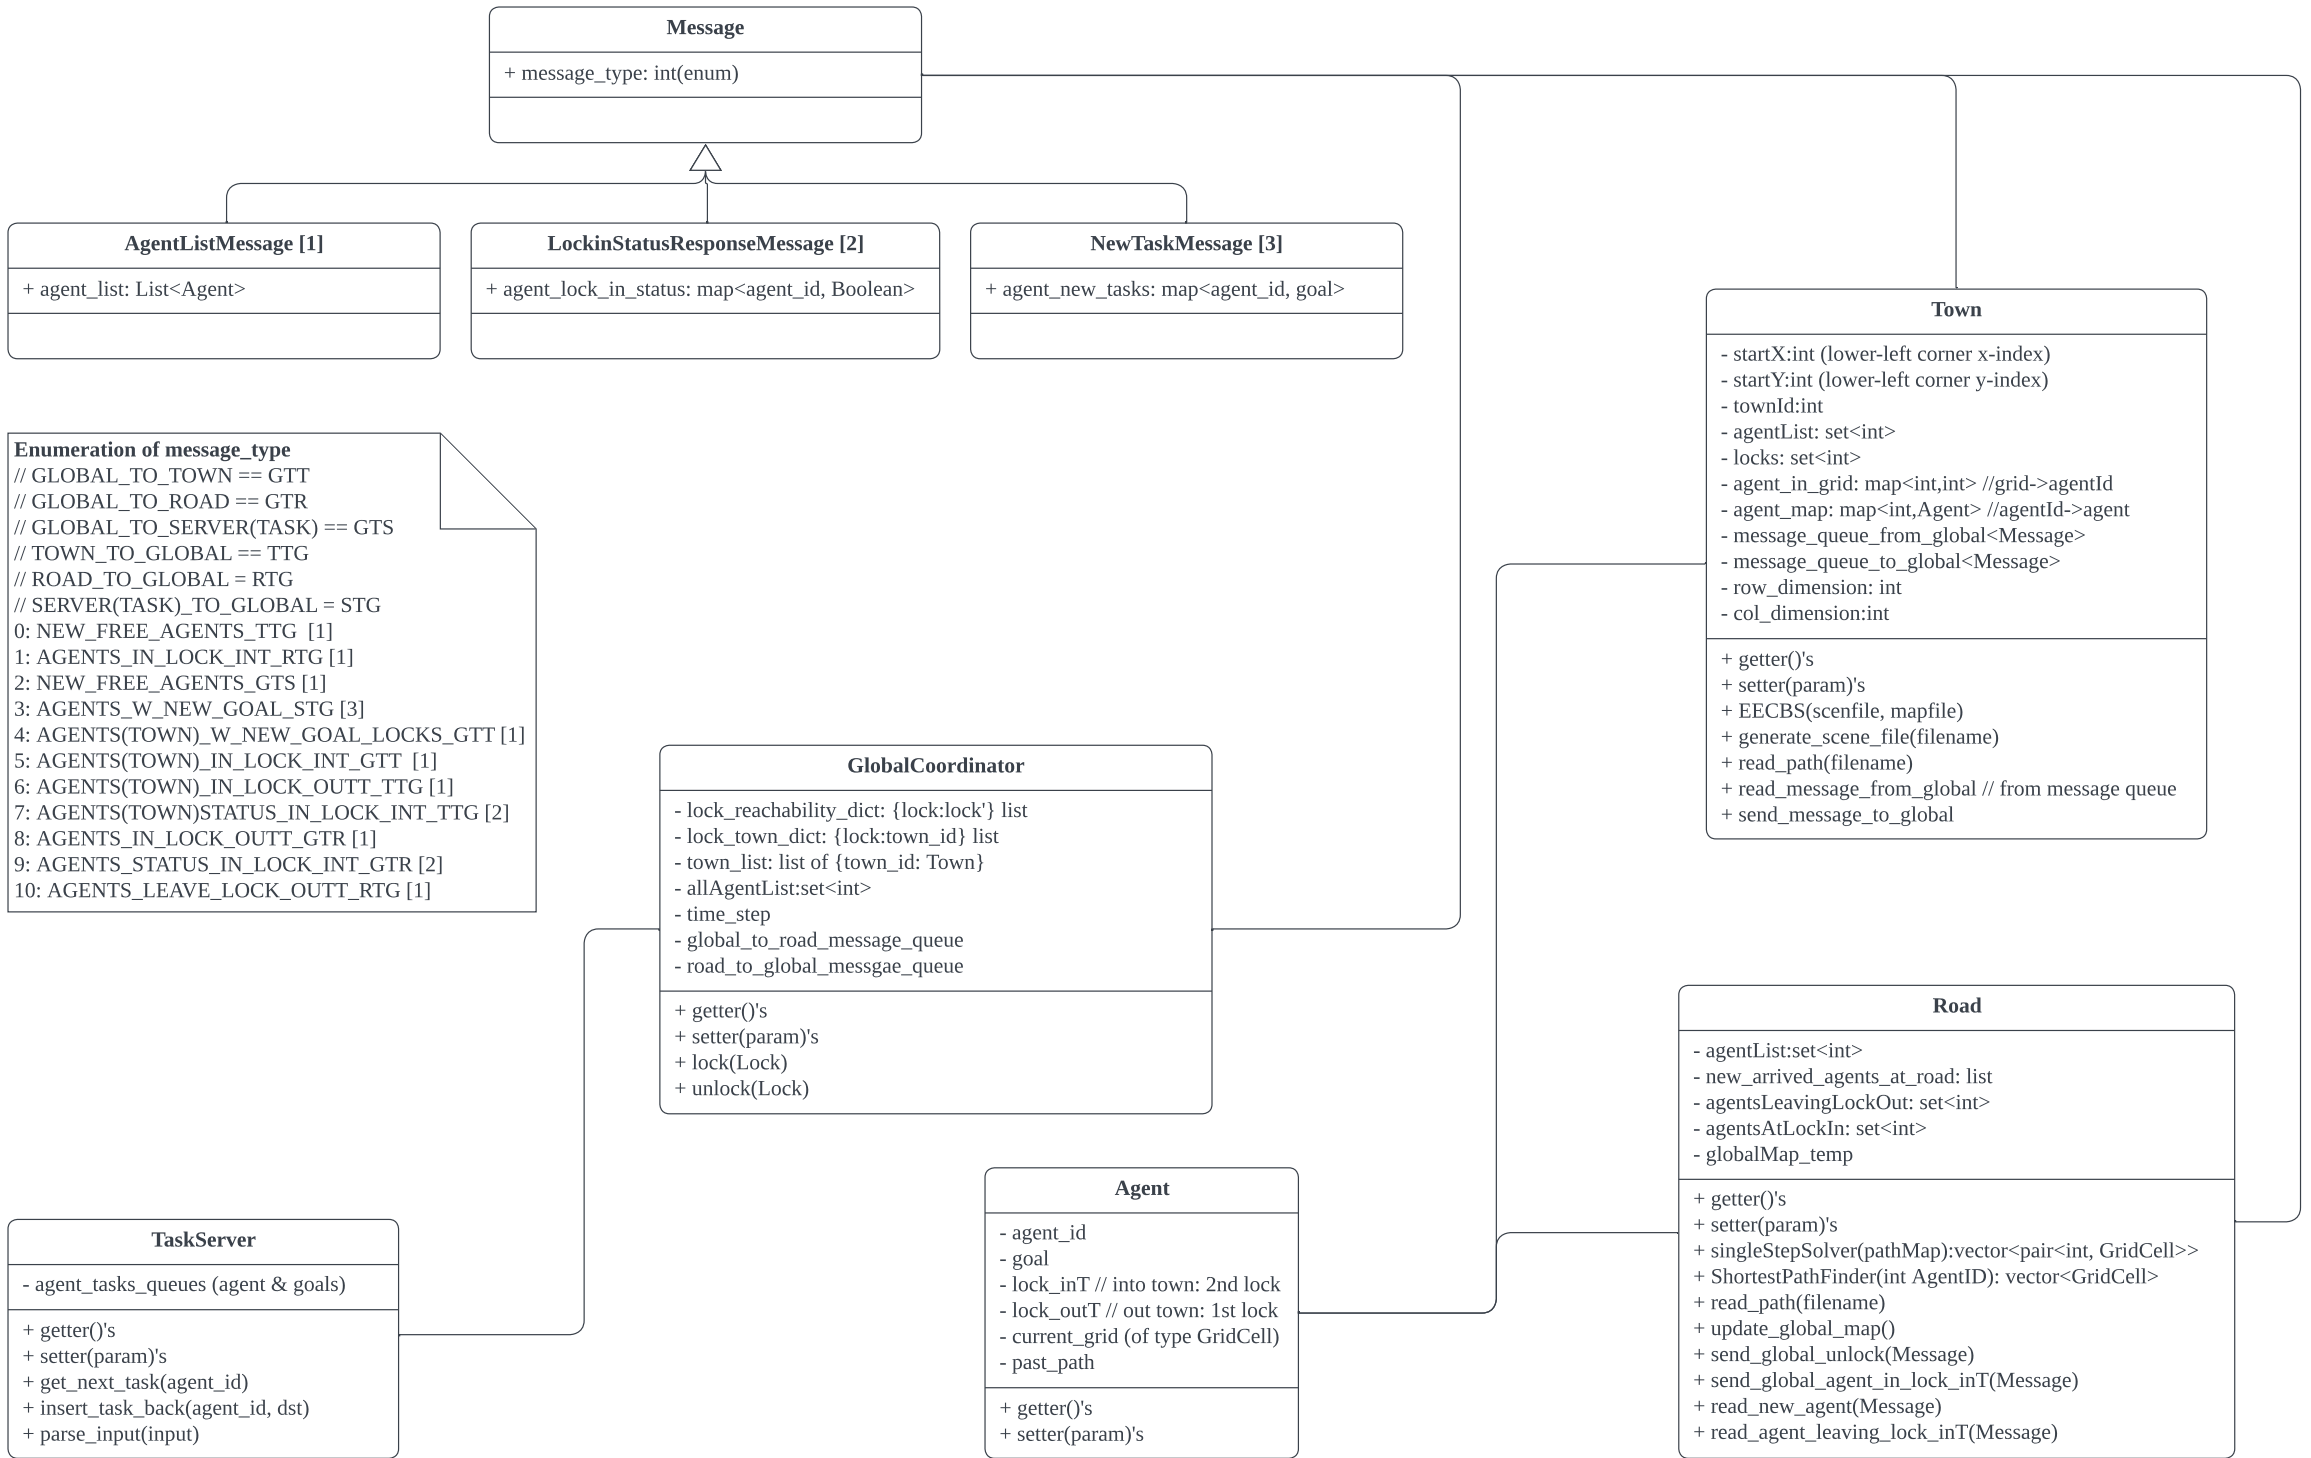
Message:

Message is used for thread communications. It contains the actual data to be sent from one component to another. For example, by the end of a timestep, each Town should construct a message that includes all new free agents that just arrived at the goal and send the message to Global. Then, the global Coordinator constructs and sends the message that assigns the goals for new free agents back to Town. The detailed message types are included in the class diagram.

5. Sequence Diagram & UML Class Diagram

Sequence diagram for One TimeStep





6. Current Main Components & Status

6a. Tigo's Town Component:

Inputs:

- Scenario files for each town(specifying agents' goal in each town)
- Map files and basic information(townID, dimension...) for each town
- All agents in the town

Outputs:

- The path for each agent in each town for each timestep from EECBS

6b. Ricardo's Highway Component (Two main functions):

1. Shortest Path Finder:

Inputs:

- The map/graph of highways
- Agent newly enters the highway from towns(also with its current and goal location, which are locks)

Outputs:

- The shortest path from the agent's current location to the goal location is a vector of (x,y) coordinates. Store the path plan for the agent.

2. Single Step Solver:

Inputs:

- All agents on highways, with their shortest paths that have been planned for them

Outputs:

- The next step of each agent. (By taking the next step from each agent's path plan, check if there is any conflict. Solve the conflict and decide the priority based on some traffic rules. If an agent can move its next step, update its path plan by deleting this step from its path plan)

6c. Dongming's Highway Layout Algorithm Component:

Input & Output of the Algorithm (written in Python):

Inputs:

- The MAPF baseline map (.map file)
 - o The whole original ascii map
 - o (.) as empty grid, (others) as obstacles
- The config file, where we specify some hyper-parameters:
 - o max_town_size: the expected maximum town size
 - o max_search_around: maximum L1-distance to search around for endpoints
 - o visited_add_cost: how much cost added to edges when lay highway
 - o max_allowed_cost: max cost allowed before blocking the edge in highway usage

Output:

- Highway File: (txt)
 - o Highway nodes & Locks as ascii map
 - o (s, n, w, e) as directions, (I) as intersection, (@) as obstacle/town, (L) as locks
- Town File: (json), for each town in the map, the file have a dict of its:
 - o Town index
 - o Town nodes & Locks as ascii map
 - o (.) as normal town node, (@) as obstacle/highway, (L) as locks
 - o Origin (left bottom corner's index)

Steps of the Algorithm:

1. G: read the original graph from input MAPF map file, only keep the largest connected component subgraph as our interest (otherwise, connectivity cannot be satisfied).
2. G_d: the directed version of G, 2 opposite direction edges for each original edge.
3. ideal_endpts: from the G's dimension and max_town_size, the ideal endpoints (not necessarily nodes) of highways so that the highways will partition the map in a way that each town has size around max_town_size.
4. actual_endpts: from ideal_endpts and max_search_around, search around the ideal endpoints for actual endpoints that should be actual nodes in the map to be highway endpoints.
5. outer_endpts_list: the outermost endpoints around the map of actual_endpts, in clockwise direction order.
6. outer_highways: plan highways between the outer_endpts_list in clockwise direction order (this aims to increase connectivity).
7. inner_highways: plan highways between each adjacent actual highway endpoints in actual_endpts.
8. Post Processing the highways and create some graphs: G_hw: graph only storing highway nodes; G_together: graph storing highway and town nodes; G_temp_locks: temporary lock locations (might be invalid); G_final_town: contains all non-highway nodes.
9. G_valid_locks: from the G_temp_locks save only the valid locks (each valid lock should be reachable from town nodes and should not disconnect the town).
10. Finally, post processing the result and create the output files.

Versions and Accomplishments in each Version:

v1, v2: not working

v3: graph edge is unweighted

- highly possible to find path with overlapping edges

v4: graph with weighted edge design

- problem: edge cases where limiting edges does not resolve conflicts at intersection
- simply removing edge/adding edge cost does not resolve node conflicts

v5: weighted & directed edge design

- when creating the original graph, each undirected edge replaced by 2 directed opposite edge (conjugate)
- when an edge is occupied by a hw, remove its conjugate => avoid actual conflict
- when an edge is occupied by a hw, add cost c in $[0, \infty)$ => avoid too much traffic
- in this way, each road can have at most 1 in/out in each direction, a max of 4 in/out

v6: add file output function

v7: add "lock" function: set lock as L in the output files

v8: what "lock"s each "lock" can reach; replace end of road with X

v9: try to fix previous lock's two problems:

1. lock not connected to any town nodes => cannot reach from town
2. lock disconnect town nodes => town become disconnected now

v10: refactor code & add documentation

v11: modifications on making the result T-connected

1. find outer highway endpoints

2. try to make them connected clockwise, before layout other part of highways
 3. * this step can be done recursively, layer by layer, possible extension...
- **Definition of T-connected:** for each town1-town2 pair, there is a direct highway from town1 to town2, where:
 - o adj 2 town nodes (within same town) are bi-direction connected
 - o adj town node & lock node are bi-direction connected
 - o adj 2 lock nodes are NOT connected
 - o adj lock node & highway node are bi-direction connected
 - o adj 2 highway nodes are only connected based on highway direction

v12: build town classes, check town connectivity

v13: add more constraint on picking outer highway endpoints

1. the end point picked is as close to the ideal as possible (max is max_search_around)
 2. the end point picked MUST be connected with 4 surrounding map nodes
- result T-connected: Boston_0_256 (0 unconnected), ost003d (0 unconnected), den312d (0 unconnected), random-64-64-20 (0 unconnected), den520d (1 unconnected)

v14: improve inner highway layout algorithm:

1. try all adjacent (s, t), same as before
2. if not found path between (s, t), try: (s-, t) & (s, t+) until find both or run out of s- & t+